


Distributed File Systems

- definition, main concepts, design goals
- semantics of file sharing
 - ◆ unix
 - ◆ session
- file access and data caching
 - ◆ cash location
 - ◆ cash modification
 - ◆ cash validation

Distributed file systems



- *Distributed file system* is a part of distributed system that provides a user with a unified view of the files on the network.

DFS - main notions


- File service — specification of the file system interface as seen by the clients
- File server — a process running on some machine which helps implement the file service by supplying files
- In principle, files in a distributed file system can be stored at any machine
 - ◆ However, a typical distributed environment has a few dedicated machines called file servers that store all the files
- A machine that holds the shared files is called a server, a machine that accesses the files is called a client.

Goals of DFS design

- Goals of a distributed file system
 - ◆ transparency
 - structure - clients should not be aware of multiple servers, replicas and caches in use
 - access - remote and local files should be accessed the same way
 - name - the name of the file should not differ on different parts of DFS
 - ◆ user mobility/file mobility - users should be able to access the DFS in a uniform manner from different location, should be able to move files around in DFS
 - ◆ simplicity/ease of use - should be similar in use to a centralized file system
 - ◆ Availability / robustness — file service should be maintained even in the presence of partial system failures
 - ◆ performance/scalability— should overcome bottlenecks of a centralized file system, should scale well

Mounting mechanism for transparency

- transparent name space can be built using Unix *mounting* mechanism
 - ◆ file systems from servers are attached (mounted) as directories in local file system
 - ◆ the points of attachment are called mount points



File sharing semantics

- Unix semantics
 - ◆ description:
 - enforces an absolute time ordering on all operations
 - every read operation on file sees the effects of all previous write operations on that file
 - ◆ can be implemented on a single-server DFS
 - ◆ easiest to use
- session semantics
 - ◆ session - series of file accesses made between open and close operations
 - ◆ changes made to the file are visible only to client process (possibly to processes on the same client)
 - ◆ the changes are visible to the sessions that open after the session closes

File access models

- Accessing remote files:
 - ◆ remote service model - client submits requests to server, all processing is done on server, file never moves from server
 - ↪ server is bottleneck
 - ↪ excessive network communication possible
 - ◆ data-caching model
 - ↪ File service provides:
 - open — transfer entire file to client
 - close — transfer entire file to server
 - ↪ Client works on file locally (in memory or on disk)
 - Simple, efficient if working on entire file
 - Must move entire file
 - Needs local disk space

Cache location

- No caching — all files on server's disk
 - ◆ Simple, no local storage needed
 - ◆ Expensive transfers
- Cache files in server's memory
 - ◆ Easy, transparent to clients
 - ◆ Still involves a network access
- Cache files on client's local disk
 - ◆ Plenty of space, reliable
 - ◆ Faster than network, slower than memory
- Cache files in client's memory
 - ◆ The usual solution (either in each process's address space, or in the kernel)
 - ◆ Fast, permits diskless workstations
 - ◆ Data may be lost in a crash

Cache modification policy (cont.)

- Variations on write-back (when are the new values flushed to the server?)
 - ◆ Write-on-close — flush new value to the server only when the file is closed
 - ↪ Can reduce disk writes, particularly when the file is open for a long time
 - ↪ Unreliable — if machine crashes, unwritten data is lost
 - ↪ May make the process wait on the file close
 - ◆ Write-periodically — flush new value to the server at periodic intervals (maybe 30 seconds)
 - ↪ Can only lose writes in last period

Remote file access and sharing

- Once the user specifies a remote file, the OS can do the access either:
 - ◆ Remotely on the server machine, and then return the results (RPC model), or
 - ◆ Can transfer the file (or part of the file) to the requesting host, and perform local accesses, or
 - ◆ Instead of doing the transfer for each user request, the OS can cache files, and use that cache to reduce the latency for data access (and thus increase performance)
- Issues
 - ◆ Where and when is data cached?
 - ◆ Cache consistency:
 - ↪ What happens when the user modifies the file? Does each cached copy change? Does the original file change?
 - ↪ Is the cached copy out of date?

Cache modification policy

- Cache modification (writing) policy decides when a modified (dirty) cache block should be flushed to the server
- Write-through — immediately flush the new value to server (& keep in cache)
 - ◆ No problems with consistency
 - ◆ Maximum reliability during crashes
 - ◆ Doesn't take advantage of caching during writes (only during reads)
- Write-back (delayed-write) — flush the new value to server after some delay
 - ◆ Fast — write need only hit the cache before the process continues
 - ◆ Can reduce disk writes since the process may repeatedly write the same location
 - ◆ Unreliable — if machine crashes, unwritten data is lost

Cache validation

- A client must decide whether or not a locally cached copy of data is consistent with the master copy
- Client-initiated validation:
 - ◆ Client initiates validity checks
 - ◆ Client contacts the server and asks if its copy is consistent with the server's copy
 - ↪ At every access, or
 - ↪ After a given interval, or
 - ↪ Only on file open
 - ◆ Server could enforce single-writer, multiple-reader semantics, but to do so
 - ↪ It would have to store client state (expensive)
 - ↪ Clients would have to specify access type (read / write) on open
 - ◆ High frequency of validity checks may mitigate the benefits of caching

Cache validation (cont.)

- Server-initiated validation:
 - ◆ Server records the parts of each file that each client caches
 - ◆ Server detects potential conflicts if two or more clients cache the same file
- ◆ Handling conflicts:
 - Session semantics — writes are only visible in sessions starting later (not to processes which have file open now)
 - When a client closes a file that it has modified, the server notifies the other clients that their cached copy is invalid, and they should discard it
 - If another client has the file open, discard it when its session is over
 - UNIX semantics — writes are immediately visible to others
 - Clients specify the type of access they want when they open a file, so if two clients want to write the same file for writing, that file is not cached
- ◆ Significant overhead at the server

Stateful vs. stateless

- Stateful server — server maintains state information for each client for each file
 - ◆ Connection-oriented (open file, read / write file, close file)
 - ◆ Enables server optimizations like read-ahead (prefetching) and file locking
 - ◆ Difficult to recover state after a crash
- Stateless server — server does not maintain state information for each client
 - ◆ Each request is self-contained (file, position, access)
 - Connectionless (open and close are implied)
 - ◆ If server crashes, client can simply keep retransmitting requests until it recovers
 - ◆ No server optimizations like above
 - ◆ File operations must be idempotent