


Distributed shared memory

- ◆ motivation and the main idea
- ◆ consistency models
 - strict and sequential
 - causal
 - PRAM and processor
 - weak and release
- ◆ implementation of sequential consistency
- ◆ implementation issues
 - granularity
 - thrashing
 - page replacement

DSM idea

- all computers share a single paged, virtual address space
- pages can be physically located on any computer
- when process accesses data in shared address space a *mapping manager* maps the request to the physical page
- mapping manager – kernel or runtime library
- if page is remote – block the process and fetch it



Advantages of DSM

- Simpler abstraction - programmer does not have to worry about data movement, may be easier to implement than RPC since the address space is the same
- easier portability - sequential programs can in principle be run directly on DSM systems
- possibly better performance
 - ◆ locality of data - data moved in large blocks which helps programs with good locality of reference
 - ◆ on-demand data movement
 - ◆ larger memory space - no need to do paging on disk
- flexible communication - no need for sender and receiver to exist, can join and leave DSM system without affecting the others
- process migration simplified - one process can easily be moved to a different machine since they all share the address space

Maintaining memory coherency

- DSM systems allow concurrent access to shared data
- concurrency may lead to unexpected results - what if the read does not return the value stored by the most recent write (write did not propagate)?
- Memory is *coherent* if the value returned by the read operation is always the value the programmer expected
- To maintain coherency of shared data a mechanism that controls (and synchronizes) memory accesses is used.
- This mechanism only allows a restricted set of memory access orderings
- *memory consistency model* - the set of allowable memory access orderings

Strict and sequential consistency

- strict consistency (strongest model)
 - ◆ value returned by a read operation is always the same as the value written by the most recent write operation
 - ◆ hard to implement
- sequential consistency (Lampert 1979)
 - ◆ the result of any execution of the operations of all processors is the same as if there were executed in some sequential order and one process' operations are in the order of the program
 - Interleaving of operations doesn't matter, if all processes see the same ordering
 - ◆ read operation may not return result of most recent write operation!
 - running a program twice may give different results
 - ◆ little concurrency

Causal consistency

- proposed (Hutto and Ahmad 1990)
- there is no single (even logical) ordering of operations – two processes may see the same operations ordered differently
- the operations are sequenced in the same order if they are potentially causally related
- read/write (or two write) operations on the same item are causally related
- all operations of the same process are causally related
- causality is transitive - if a process carries out an operation B that causally depends on the preceding op A - all consequent ops by this process are causally related to A (even if they are on different items)

PRAM and processor consistency

- PRAM (Lipton & Sandberg 1988)
 - ◆ All processes see only memory writes done by a single process in the same (correct) order
 - ◆ PRAM = pipelined RAM
 - ↗ Writes done by a single process can be pipelined; it doesn't have to wait for one to finish before starting another
 - ↗ writes by different processes may be seen in different order on a third process
 - ◆ Easy to implement — order writes on each processor independent of all others
- Processor consistency (Goodman 1989)
 - ◆ PRAM +
 - ◆ coherency on the same data item - all processes agree on the order of write operations to the same data item

Comparison of consistency models

- Models differ by difficulty of implementation, ease of use, and performance
- Strict consistency — most restrictive, but hard to implement
- Sequential consistency — widely used, intuitive semantics, not much extra burden on programmer
 - ◆ But does not allow much concurrency
- Causal & PRAM consistency — allow more concurrency, but have non-intuitive semantics, and put more of a burden on the programmer to avoid doing things that require more consistency
- Weak and Release consistency — intuitive semantics, but put extra burden on the programmer

Implementing sequential consistency on page-based DSM

- Can a page move? ...be replicated?
- Nonreplicated, nonmigrating pages
 - ◆ All requests for the page have to be sent to the owner of the page
 - ◆ Easy to enforce sequential consistency — owner orders all access request
 - ◆ No concurrency
- Nonreplicated, migrating pages
 - ◆ All requests for the page have to be sent to the owner of the page
 - ◆ Each time a remote page is accessed, it migrates to the processor that accessed it
 - ◆ Easy to enforce sequential consistency — only processes on that processor can access the page
 - ◆ No concurrency

Weak and release consistency

- Weak consistency (Dubois 1988)
 - ◆ Consistency need only apply to a group of memory accesses rather than individual memory accesses
 - ◆ Use synchronization variables to make all memory changes visible to all other processes (e.g., exiting critical section)
 - ↗ all access to synchronization variables must be sequentially consistent
 - ↗ write operations are completed before access to synchvar
 - ↗ access to non-synchvar is allowed only after synchvar access is completed
- Release consistency (Gharachorloo 1990)
 - ◆ two synchronization vars
 - ↗ acquire - all changes to synchronized vars are propagated to the process
 - ↗ release - all changes to synchronized vars are propagated to other processes
 - ↗ programmer has to write accesses to these variables

Implementation issues

- how to keep track of the location of remote data
- how to overcome the communication delays and high overhead associated with execution of communication protocols
- how to make shared data concurrently accessible at several nodes to improve system performance

Implementing sequential consistency on page-based DSM (cont.)

- Replicated, migrating pages
 - ◆ All requests for the page have to be sent to the owner of the page
 - ◆ Each time a remote page is accessed, it's copied to the processor that accessed it
 - ◆ Multiple read operations can be done concurrently
 - ◆ Hard to enforce sequential consistency — must invalidate (most common approach) or update other copies of the page during a write operation
- Replicated, nonmigrating pages
 - ◆ Replicated at fixed locations
 - ◆ All requests to the page have to be sent to one of the owners of the page
 - ◆ Hard to enforce sequential consistency — must update other copies of the page during a write operation

Granularity

- Granularity - size of shared memory unit
- Page-based DSM
 - ◆ Single page — simple to implement
 - ◆ Multiple pages — take advantage of locality of reference, amortize network overhead over multiple pages
 - ↳ Disadvantage — false sharing
- Shared-variable DSM
 - ◆ Share only those variables that are needed by multiple processes
 - ◆ Updating is easier, can avoid false sharing, but puts more burden on the programmer
- Object-based DSM
 - ◆ Retrieve not only data, but entire object — data, methods, etc.
 - ◆ Have to heavily modify old programs

Thrashing

- Occurs when system spends a large amount of time transferring shared data blocks from one node to another (compared to time spent on useful computation)
 - ◆ interleaved data access by two nodes causes data block to move back and forth
 - ◆ read-only blocks invalidated as soon as they are replicated
- handling thrashing
 - ◆ application specifies when to prevent other nodes from moving block - has to modify application
 - ◆ "nailing" block after transfer for a minimum amount of time t - hard to select t , wrong selection makes inefficient use of DSM
 - ↳ adaptive nailing?
 - ◆ tailoring coherence semantics (Minin) to use — object based sharing

Page replacement

- What to do when local memory is full?
 - ◆ swap on disk?
 - ◆ swap over network?
 - ◆ what if page is replicated?
 - ◆ what if it's read-only?
 - ◆ what if it's read/write but clean(dirty)?
 - ◆ are shared pages given priority over private (non-shared)?