

Fault tolerance in distributed systems

- Motivation
- robust and stabilizing algorithms
- failure models
- robust algorithms
 - ◆ decision problems
 - ◆ impossibility of consensus in asynchronous networks with crash-failures
 - ◆ consensus and agreement with initially-dead processes - knot calculation algorithm
- stabilization
 - ◆ Dijkstra's K-state algorithm

1

Why fault tolerance

- Distributed systems encompass more and more individual devices
- the chance of failure in distributed system can grow arbitrarily large when the number of its components increases
- distributed systems can hardly be restarted after failure
- distributed systems are subjects to *partial failure* property: when one of the components fails the system may still be able to function in a decreased capacity
- as the system grows in size
 - ◆ it becomes more likely that one component fail
 - ◆ it becomes less likely that the failure occurs in all components
- thus the systems able to deal with failures are attractive

2

Robust and stabilizing algorithms

- An algorithm is *robust (masking)* if the correct operation of the algorithm is ensured even at the presence of specified failures
- the algorithm is stabilizing if it is able to eventually start working correctly regardless of the initial state.
 - ◆ stabilizing algorithm does not guarantee correct behavior during recovery
 - ◆ stabilizing algorithm is able to recover from faults regardless of their nature (as soon as the influence of the failure stops)
- an algorithm can mask certain kinds of failures and stabilize from others
 - ◆ for example: an algorithm may mask message loss and stabilize from topology changes

3

Failure Models

- Faults form a hierarchy on the basis of the severity of faults
- benign
 - ◆ initially dead - a process is initially dead if it does not execute a single step in its algorithm
 - ◆ crash model - a process executes steps correctly up to some moment (crash) and stops executing thereafter
- malign - Byzantine - a process executes arbitrary steps (not necessarily in accordance with its local algorithm). In particular Byzantine process sends messages with arbitrary content
- initially dead process is a specially case crashed process which is a special case of Byzantine process
 - ◆ if algorithm is Byzantine-robust it can also tolerate crashes and initially dead processes
 - ◆ if a problem cannot be solved for initially dead processes, it cannot be solved in the presence of crashes or Byzantine failures
- other fault models can be defined in between

4

Decision problems

- Study of robust algorithms centers around decision problems
- *decision problem* requires that each (correct) process eventually and irreversibly arrives at a "decision" value
- decision problems requirements:
 - ◆ termination - all correct processes decide (cannot indefinitely wait for dead processes)
 - ◆ consistency - the decisions of correct processes should be related;
 - ↳ *consensus problem* - the decisions are equal
 - ↳ *election problem* - only one process arrives at "1" (leader) the others - "0" (non-leaders)
 - ◆ non-triviality - different outputs are possible in different executions of the algorithm

5

Impossibility of consensus

- State is *reachable* if there is a computation that contains it
- Each process has a read-only input variable x_p and write-once output variable y_p initially holding b
- A consensus algorithm is 1-crash robust if it satisfies the following properties:
 - ◆ termination - in every 1-crash fair execution all correct processes decide
 - ◆ agreement - if, in any reachable state, $y_p \neq b$ and $y_q \neq b$ for correct processes p and q then, $y_p = y_q$
 - ◆ non-triviality - there exist a reachable states such that for some p , $y_p=1$ in one state and $y_p=0$ in another
- Theorem: there are no asynchronous, deterministic 1-crash robust consensus algorithms
- intuitively this result is explained by the fact that in an asynchronous system it is impossible to distinguish a crashed process from an infinitely slow one

6

What is possible

- Consensus with initially dead-process fault model is possible
- weaker coordination problems than consensus (such as renaming) are solvable
 - given: a set of processes p_1, \dots, p_N , each process with distinct identity taken from arbitrary large domain. Each process has to decide on a unique new name from smaller domain $1, \dots, K$
- randomized algorithms are possible even for Byzantine failures
- weak termination - termination required only when a given process (general) is correct, the objective is for all processes to learn the general's decision; solvable even in the presence of Byzantine faults
- synchronous systems are significantly more fault tolerant

7

Consensus with initially dead processes

- If processes are only initially-dead consensus is possible.
- Based on the following knot-computation algorithm
- knot is a strongly connected sub-graph with no outgoing edges
- the objective is for all correct processes to agree on the subset of correct processes
- L stands for $\lfloor (N+1)/2 \rfloor$
- we assume that there are at least L alive processes
- first phase: each process p :
 - sends messages to all processes in the system
 - collects at least L messages in set $Succ_p$
- a process is a successor if p got a message from it - there is a graph G in the system
- thus each correct process has L successors
- an initially-dead process does not send any messages. Thus there is a knot in G containing correct processes

8

Knot calculation algorithm

```

var  $Succ_p, Alive_p, Rcvd_p$  : sets of processes    init  $\emptyset$  ;
begin shout (name, p) ;
  (* that is: forall  $q \in \mathbb{P}$  do send (name, p) to q *)
  while # $Succ_p < L$ 
    do begin receive (name, q) ;  $Succ_p := Succ_p \cup \{q\}$  end ;
  shout (pre, p,  $Succ_p$ ) ;
   $Alive_p := Succ_p$  ;
  while  $Alive_p \not\subseteq Rcvd_p$ 
    do begin receive (pre, q,  $Succ$ ) ;
         $Alive_p := Alive_p \cup Succ \cup \{q\}$  ;
         $Rcvd_p := Rcvd_p \cup \{q\}$ 
      end ;
  Compute a knot in  $G$ 
end
    
```

9

Knot calculation algorithm (cont.)

- Since each correct process has an outdegree L - the knot has at least L processes
- since $L > N/2$, G contains just one knot. Let's call it K
- since p has L successors, one of them is in K , thus all nodes in K are descendants of p
- second phase:
 - each process collects a list of its descendants. Since processes do not fail, no deadlock occurs at this stage
- in the end a process has a set of processes and their descendants which allows it to compute the knot in G
- it is possible to do an election and consensus on the basis of knot calculation algorithm
 - election - since processes agree on the knot - they all can agree on the leader by electing the leader - the process with the highest id in K
 - consensus - all processes calculate the input value of the processes of K and output the value that occurs most often

10

Guarded Command Language (GCL)

```

*[
  guard1 → command1
  [] guard2 → command2
  :
  ]
    
```

- *[...] - execution repeats forever
 - guard _{i} - binary predicate on local vars, received messages, etc.;
 - command _{i} - list of assignment statements;
- command is executed when corresponding guard is true; guards are selected nondeterministically,

Advantages:

- GCL allows to easily reason about algorithms and their executions: the program counter position is irrelevant or less important;
- we don't have to consider execution starting in the middle of guard or command (*serializability property*);

11

Dijkstra's K-State Token Circulation Algorithm

Objective: circulate a single token among processors

```

Processor  $p^0$ 
*[
   $s^0 = s^{k-1} \rightarrow s^0 := (s^0 + 1) \bmod K$ 
,

Processor  $p^i$  ( $0 < i < K$ )
*[
   $s^i \neq s^{i-1} \rightarrow s^i := s^{i-1}$ 
,
    
```

- the system consists of a ring of K processors (*ids* 0 through $K-1$)
- each processor maintains a state variable s_i ; a processor can see the state of its left (smaller id) neighbor
- guard evaluates to **true** - processor has a privilege (token)
- all processors evaluate their guards, only **one at a time** changes state (C-Daemon)
- after the state change all processors re-evaluate the guards

12