

MPP vs Hadoop

Alexey Grishchenko

HUG Meetup
28.11.2015

Agenda

- Distributed Systems
- MPP
- Hadoop
- MPP vs Hadoop
- Summary

Agenda

- **Distributed Systems**
- MPP
- Hadoop
- MPP vs Hadoop
- Summary

Distributed Systems


Avoid distributed systems in all the problems that potentially could be solved using non-distributed systems

Distributed Systems

- Consensus problem
 - Paxos
 - RAFT
 - ZAB
 - etc.
- Transaction consistency
 - 2PC
 - 3PC

Distributed Systems

- CAP Theorem



Distributed Systems

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

<http://www.cs.cornell.edu/projects/ladis2009/talks/dean-keynote-ladis2009.pdf>

Distributed Systems

Reasons to use

- Performance issues
 - More than 100'000 TPS
 - More than 4 GB/sec scan rate
 - More than 100'000 IOPS
- Capacity issues
 - More than 50TB of data
- DR and Geo-Distribution

Agenda

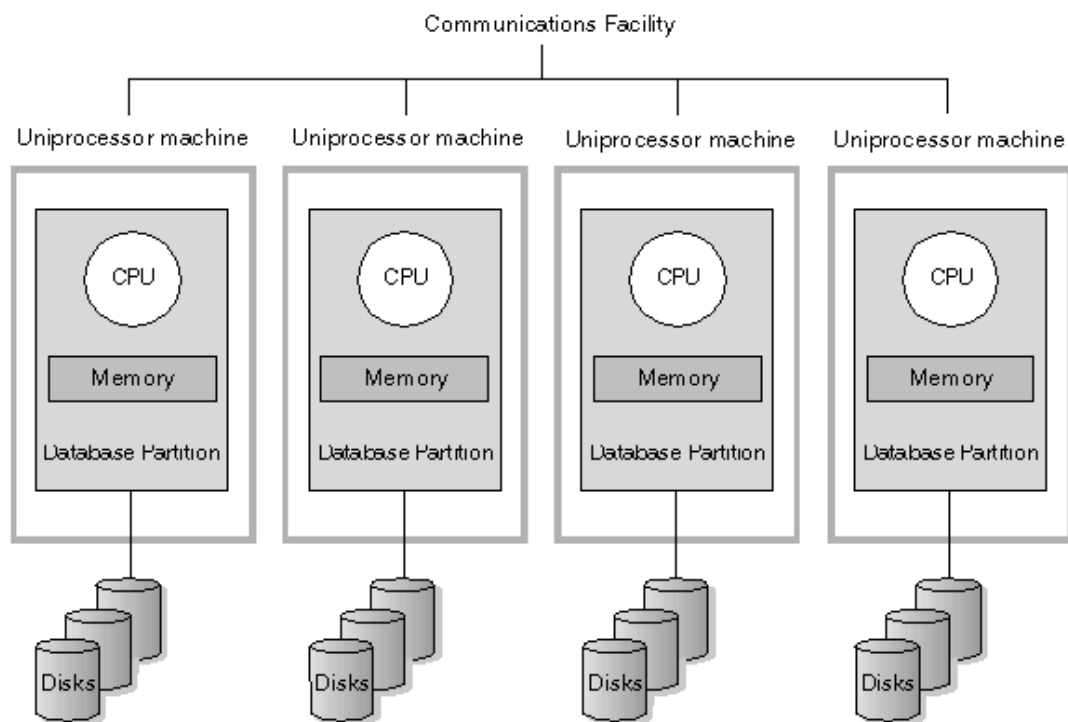
- *Distributed Systems*
- **MPP**
- Hadoop
- MPP vs Hadoop
- Summary

MPP

Main principles

- Shared Nothing
- Data Sharding
- Data Replication
- Distributed Transactions
- Parallel Processing

MPP



MPP

Works well for

- Relational data
- Batch processing
- Ad hoc analytical SQL
- Low concurrency
- Applications requiring ANSI SQL

MPP

Not the best choice for

- Non-relational data
- OLTP and event stream processing
- High concurrency
- 100+ server clusters
- Non-analytical use cases
- Geo-Distributed use cases

Agenda

- *Distributed Systems*
- *MPP*
- **Hadoop**
- MPP vs Hadoop
- Summary

Hadoop

Main Components

- HDFS
- YARN
- MapReduce
- HBase
- Hive / Hive+Tez

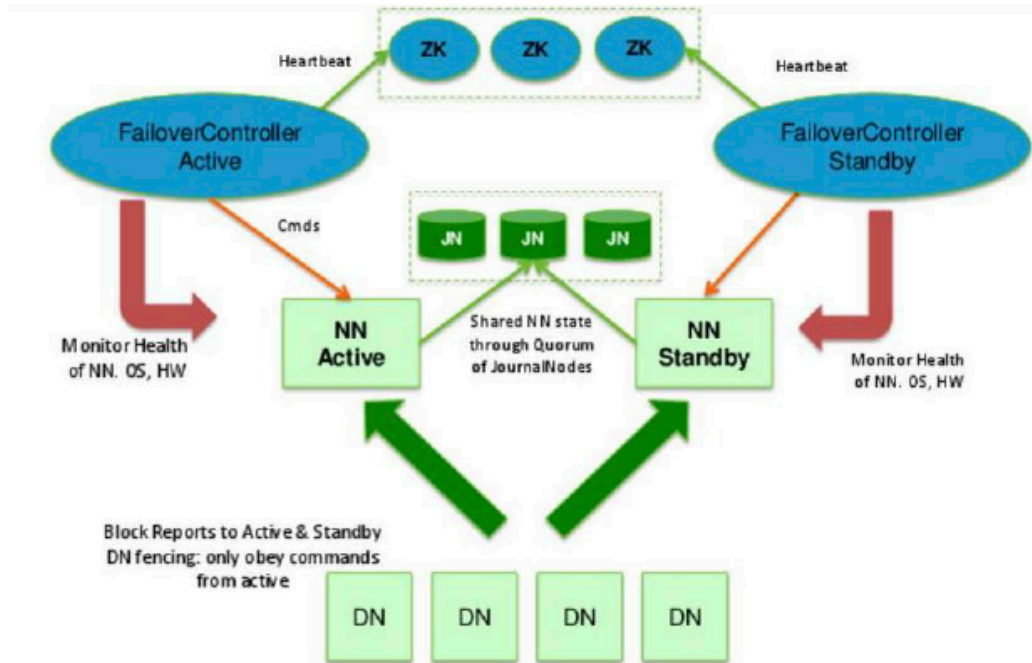
Hadoop

HDFS

- Distributed filesystem
- Block-level storage with big blocks
- Non-updatable
- Synchronous block replication
- No built-in Geo-Distribution support
- No built-in DR solution

Hadoop

HDFS



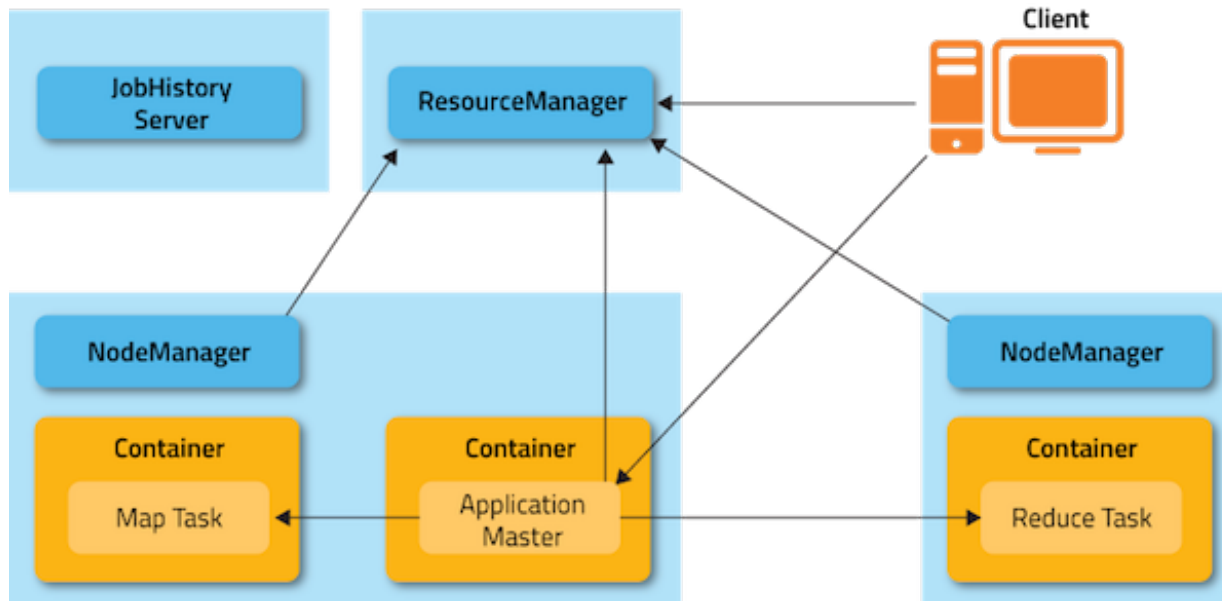
Hadoop

YARN

- Cluster resource manager
- Manages CPU and RAM allocation
- Schedulers are pluggable
- Can handle different resource pools
- Supports both MR and non-MR workload

Hadoop

YARN



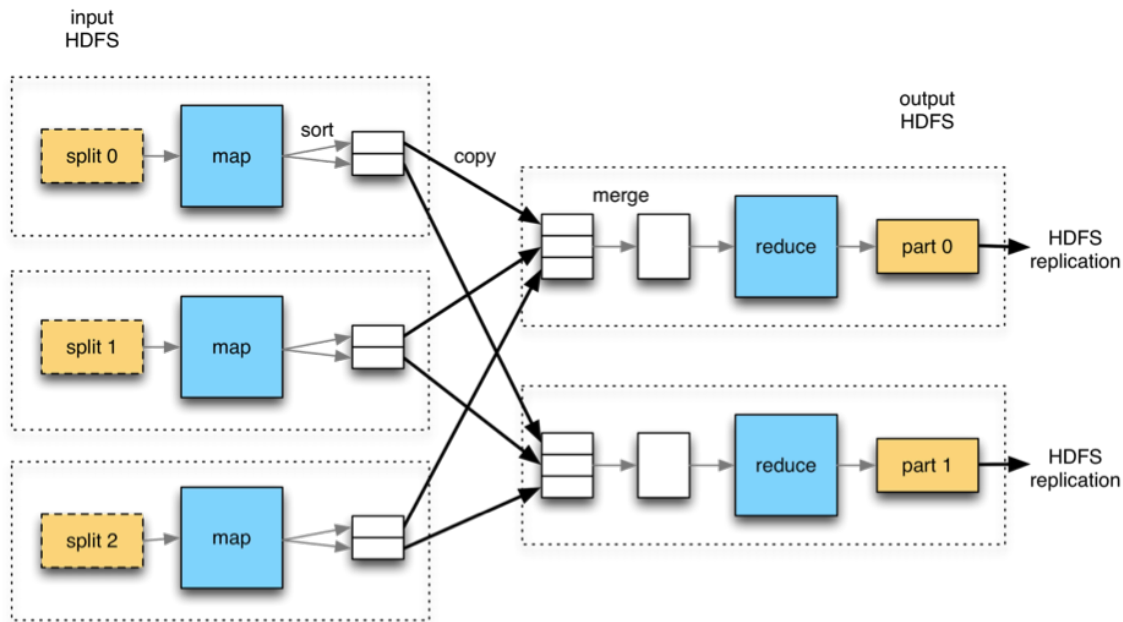
Hadoop

MapReduce

- Framework for distributed data processing
- Two main operations: map and reduce
- Data hits disk after “map” and before “reduce”
- Scales to thousands of servers
- Can process petabytes of data
- Extremely reliable

Hadoop

MapReduce



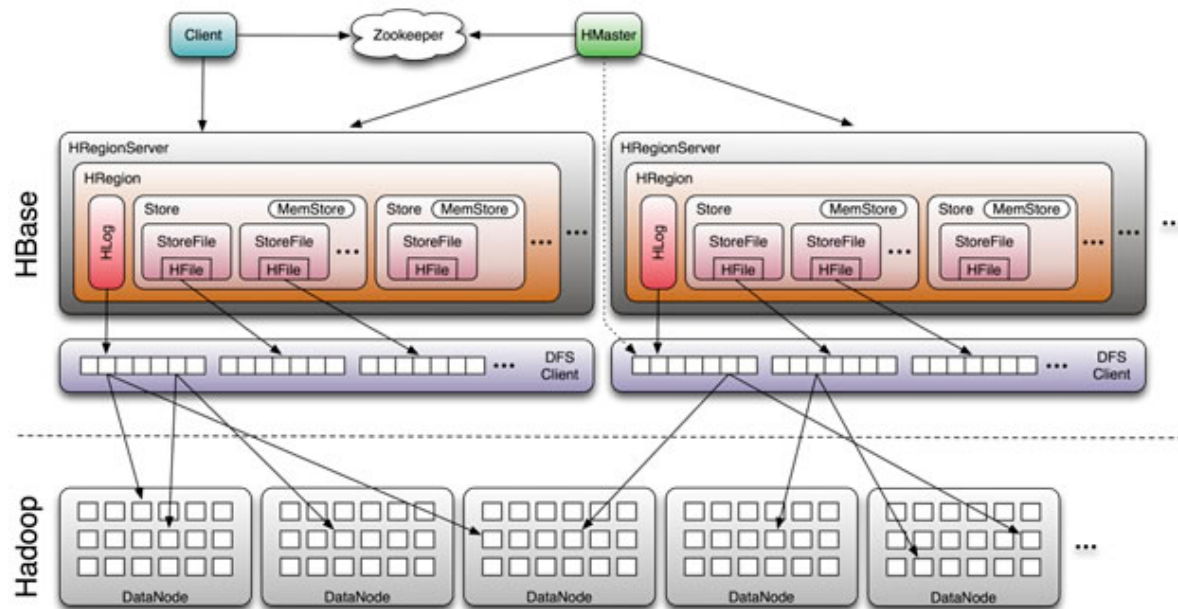
Hadoop

HBase

- Distributed key-value store
- Data is sharded by key
- Data is stored in sorted order
- Stores multiple versions of the row
- Easily scales

Hadoop

HBase



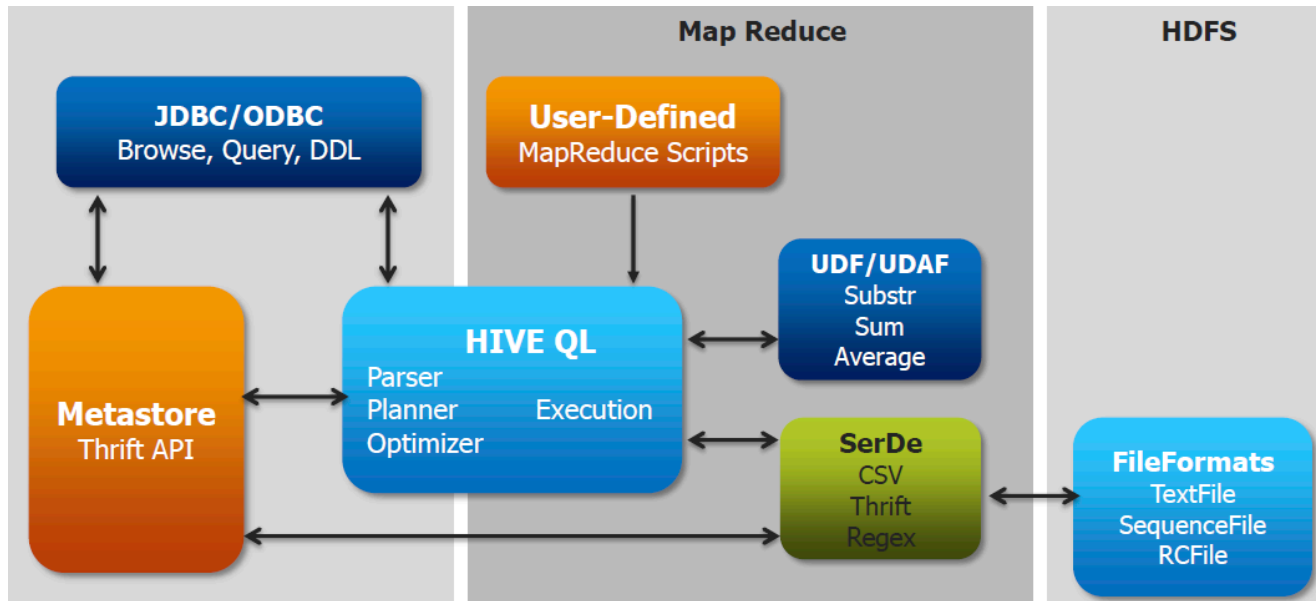
Hadoop

Hive

- Query engine with SQL-like syntax
- Translates HiveQL query to MR / Tez / Spark job
- Processes HDFS data
- Supports UDFs and UDAFs

Hadoop

Hive



Hadoop

Works well for

- Write Once Read Many
- 100+ server clusters
- Both relational and non-relational data
- High concurrency
- Batch processing and analytical workload
- Elastic scalability

Hadoop

Not the best choice for

- Write-heavy workloads
- Small clusters
- Analytical DWH cases
- OLTP and event stream processing
- Cost savings

Agenda

- *Distributed Systems*
- *MPP*
- *Hadoop*
- **MPP vs Hadoop**
- Summary

MPP vs Hadoop for Business

MPP

Hadoop

Platform Openness

Mostly Closed

Open

MPP vs Hadoop for Business

MPP

Platform Openness

Mostly Closed

Hardware Options

Mostly Appliances

Hadoop

Open

Commodity

MPP vs Hadoop for Business

MPP

Platform Openness

Hardware Options

Vendor Lock-in

Mostly Closed

Mostly Appliances

Typical

Hadoop

Open

Commodity

Not Common

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source
Supportability	Easy	Complex

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source
Supportability	Easy	Complex
Scalability	Up to 100 servers	Up to 5000 servers

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source
Supportability	Easy	Complex
Scalability	Up to 100 servers	Up to 5000 servers
Scalability	Up to 100-300 TB	Up to 100 PB

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source
Supportability	Easy	Complex
Scalability	Up to 100 servers	Up to 5000 servers
Scalability	Up to 100-300 TB	Up to 100 PB
Target Systems	DWH	Purpose-Built Batch

MPP vs Hadoop for Business

	MPP	Hadoop
Platform Openness	Mostly Closed	Open
Hardware Options	Mostly Appliances	Commodity
Vendor Lock-in	Typical	Not Common
Technology Price	\$200K – \$10M	\$50K – \$500K
Implementation Cost	Moderate	High
Extensibility	Vendor-provided APIs	Open Source
Supportability	Easy	Complex
Scalability	Up to 100 servers	Up to 5000 servers
Scalability	Up to 100-300 TB	Up to 100 PB
Target Systems	DWH	Purpose-Built Batch
Target End Users	Business Analysts	Developers

MPP vs Hadoop for Architect

Query Optimization

MPP

Good

Hadoop

Poor to None

MPP vs Hadoop for Architect

Query Optimization

MPP

Good

Hadoop

Poor to None

Debugging

Easy

Very Hard

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High
Query Latency	10-20 ms	10-20 sec

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High
Query Latency	10-20 ms	10-20 sec
Query Runtime	5-7 sec	10-15 mins

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High
Query Latency	10-20 ms	10-20 sec
Query Runtime	5-7 sec	10-15 mins
Query Max Runtime	1-2 hours	1-2 weeks

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High
Query Latency	10-20 ms	10-20 sec
Query Runtime	5-7 sec	10-15 mins
Query Max Runtime	1-2 hours	1-2 weeks
Min Collection Size	Megabytes	Gigabytes

MPP vs Hadoop for Architect

	MPP	Hadoop
Query Optimization	Good	Poor to None
Debugging	Easy	Very Hard
Accessibility	SQL	Mainly Java
DBA Skill Level	Low	High
Single Job Redundancy	Low	High
Query Latency	10-20 ms	10-20 sec
Query Runtime	5-7 sec	10-15 mins
Query Max Runtime	1-2 hours	1-2 weeks
Min Collection Size	Megabytes	Gigabytes
Max Concurrency	10-15 queries	70-100 jobs

Agenda

- *Distributed Systems*
- *MPP*
- *Hadoop*
- *MPP vs Hadoop*
- *Examples*
- **Summary**

Summary

Use MPP for

- Analytical DWH
- Ad hoc analyst SQL queries and BI
- Keep under 100TB of data

Use Hadoop for

- Specialized data processing systems
- Over 100TB of data

Questions?

Questions?

Pivotal

BUILT FOR THE SPEED OF BUSINESS